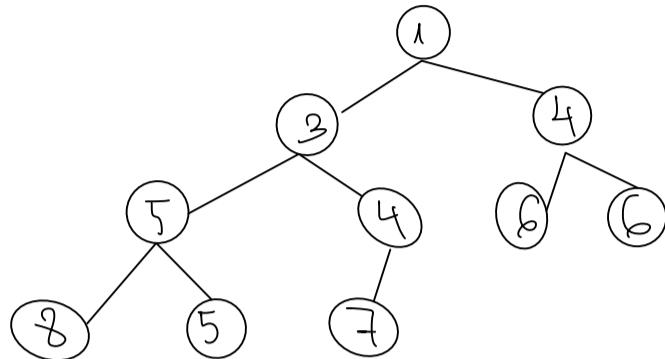


FIBONACCI HEAPS

Recall Binary Heap:

- A binary heap, satisfies these 2 properties:
- Every level is full, last one from left to right.
- Heap property: No child is smaller than its parent



Costs of basic operations in binary heap:

- INSERT : $\Theta(\lg n)$
- MINIMUM: $\Theta(1)$
- EXTRACT-MIN: $\Theta(\lg n)$
- UNION : $\Theta(n)$
- DECREASE-KEY: $\Theta(\lg n)$
- DELETE: $\Theta(\lg n)$

Observations: Most of the time, binary heaps work well. However, if we need to use UNION operation or we need constant performance for all operations for INSERT and DECREASE-KEY operations, then Fibonacci Heap works better.

Costs of Fibonacci Heap operations:

- INSERT : $\Theta(1)$
- MINIMUM: $\Theta(1)$
- EXTRACT-MIN: $\Theta(\lg n)$
- UNION : $\Theta(1)$
- DECREASE-KEY: $\Theta(1)$
- DELETE: $\Theta(\lg n)$

\Rightarrow Goal of Fibonacci Heap:

- Allow us to be lazy, i.e. do as little work as possible (INSERT, UNION, DECREASE-KEY, ...) and only performs when necessary (EXTRACT-MIN)

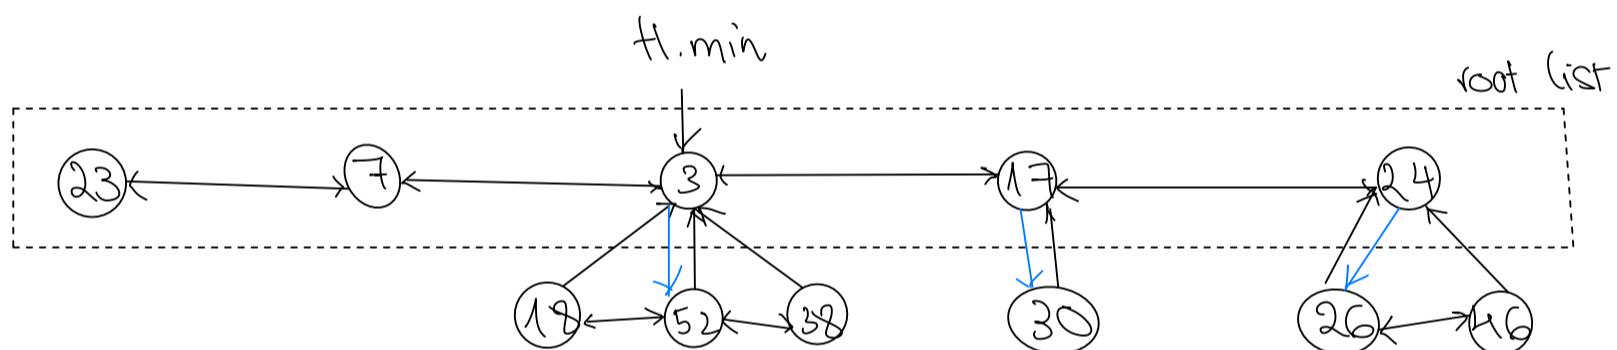
Structure of Fibonacci Heap:

A node in Fibonacci heap is defined with these properties

Heap H :	$H.\text{min}$ point to min	Node x :	$x.\text{parent}$ point to parent node
	$H.n$ number nodes in heap	$x.\text{child}$ point to any of the children	
		$x.\text{left}$	
		$x.\text{right}$	doubly linked list
		$x.\text{degree}$ number of children	
		$x.\text{mark}$ lost child's indicator:	

- marked when node is a child
- unmarked when node has no child

- Fibonacci Heap has many roots that are connected together through "left" and "right" pointers as doubly linked list. All children in the same level of a root are also connected in similar manner.



Amortized Analysis Fibonacci Heap

◦ Potential Function:

The potential function used to analyze amortized cost of Fibonacci Heap is defined as:

$$\phi(H) = \underbrace{t(H)}_{\# \text{ trees}} + 2 \cdot \underbrace{m(H)}_{\# \text{ marked nodes}}$$

Assuming we start with empty heap, the initial potential is 0. And the potential remain non negative throughout the operations

INSERT (H, c)

- The philosophy of Fibonacci heap is to delay the expensive work for as long as possible.
- In the case of insert, it simply add new node into the root list.
- The change in potential is therefore:

$$\begin{aligned}\Delta \phi &= \phi(H') - \phi(H) \\ &= [t(H') + 2 \cdot m(H')] - [t(H) + 2 \cdot m(H)] \\ &= t(H) + 1 + 2m(H) - t(H) - 2m(H) \\ &= 1\end{aligned}$$

The amortized cost:

$$\begin{aligned}\hat{c} &= c + \Delta \phi \\ &= O(1) + 1 \\ &= O(1)\end{aligned}$$

UNION (H_1, H_2)

- This operation unites H_1 and H_2 , destroy both in the process. It concatenates H_1 and H_2 based on the root with smaller value.
 - The change in potential is:
- $$\begin{aligned}\Delta \phi &= \phi(H_{12}) - (\phi(H_1) + \phi(H_2)) \\ &= (t(H_{12}) - 2m(H_{12})) - (t(H_1) + 2m(H_1) + t(H_2) + 2m(H_2)) \\ &= 0 \quad \text{because } \begin{cases} t(H_{12}) = t(H_1) + t(H_2) \\ m(H_{12}) = m(H_1) + m(H_2) \end{cases}\end{aligned}$$

The amortized cost is:

$$\begin{aligned}\hat{c} &= c + \Delta \phi \\ &= O(1) + 0 \\ &= O(1)\end{aligned}$$

EXTRACT-MIN (H)

- This is where the delayed work is carried out, it consolidate trees in the root list.

Pseudo code:

```
EXTRACT-MIN ( $H$ ) :  
    min-node =  $H.\text{min}$   
    if min-node = NIL:  
        return min-node } empty heap  
  
    for child : min-node.children:  
        add child to root-list  
        child.parent := NIL  
    remove min-node from root-list  
  
    if min-node = min-node.right:  
         $H.\text{min} := \text{NIL}$   
    else:  
         $H.\text{min} := \text{min-node.right}$   
        CONSOLIDATE ( $H$ )  
  
 $H.n -= 1$   
return min-node
```

move children to root list,
remove minimum node

min node is the last node in
root list

set new min node to any neighbor
then consolidate

CONSOLIDATE (H)

```
init array  $A[0 \dots D(H.n)]$  to nil  
for each node  $w$  in root-list:  
     $x, d := w, w.\text{degree}$   
    while  $A[d] \neq \text{NIL}$ :  
         $y := A[d]$   
        if  $x.\text{key} < y.\text{key}$ : swap  $x, y$   
        FIB-HEAP-LINK ( $H, y, x$ )
```

ensure roots have
unique degree, by
merging those with
the same degree

```
 $A[d] := \text{NIL}$   
 $d += 1$   
 $A[d] := x$ 
```

```
 $H.\text{min} := \text{NIL}$   
for  $d : 0 \rightarrow D(H.n)$ :  
    if  $A[d] = \text{NIL}$ : continue  
    create/append to root list, update  $H.\text{min}$ 
```

build new heap
from array A
of degree .

FIB-HEAP LINK(H, y, x)

remove y from root list

make y a child of x , increase $x.\text{degree}$

$y.\text{mark} := \text{FALSE}$

- The change in potential is defined as:

$$\begin{aligned}
 \Delta \phi &= \phi(H') - \phi(H) && H': \text{heap after extract-min} \\
 &= t(H') + 2m(H') - t(H) - 2m(H) \\
 &= (D(n) + 1) + 2m(H) - t(H) - 2m(H) \\
 &\quad \text{unique roots} \qquad \text{no addition} \\
 &= D(n) + 1 - t(H) \quad \sim \quad O(D(n) - t(H))
 \end{aligned}$$

- A amortized cost:

$$\begin{aligned}
 \hat{c} &= c + \Delta \phi \\
 &= O(D(n) + t(H)) + O(D(n) - t(H)) \\
 &= O(2D(n)) \sim O(D(n))
 \end{aligned}$$

DECREASE_KEY (H, x, k)

Pseudo Code:

```

DECREASE_KEY(H, x, k)
if xc > x.key:
    throw Error
xc.key := k
y := xc.parent
if y != NIL and xc.key < y.key
    CUT(H, xc, y)
    CASCADING-CUT(H, y)
if xc.key < H.min.key:
    H.min = xc

```

CUT (H, xc, y)

remove xc from child list of y, decrementing y.degree
add xc to the root list of H
xc.parent := NIL
xc.mark := FALSE

CASCADING-CUT (H, y)

```

z := y.parent
if z == NIL: return
if y.mark == FALSE:
    y.mark := TRUE
return

```

base case

CUT (H, y, z)

CASCADING-CUT (H, z)

The actual cost: CUT takes only $O(1)$, but it is a recursive func that goes up the tree and cut any node that loses its second child. So: $O(c)$, where c is the number of calls to the CASCADING-CUT func.

• The potential change:

$$\begin{aligned}
 \Delta \phi &= \phi(H') - \phi(H) \\
 &= t(H') + 2m(H') - t(H) - 2m(H) \\
 &= (\cancel{t(H)} + c) + 2(m(H) - \cancel{(c-2)}) - t(H) - 2m(H) \\
 &\quad \text{because } c \text{ new trees created} \qquad \text{because } c \text{ cuts means } (c-1) \text{ cascading} \\
 &\Rightarrow (c-1) \text{ unmarks} \\
 &\quad - 1 \text{ potential mark at the end} \\
 &\Rightarrow c-1-1 = c-2 \\
 &= 4 - c
 \end{aligned}$$

• Amortized cost:

$$\begin{aligned}
 \hat{c} &= c + \Delta \phi \\
 &= O(c) + 4 - c \\
 &= O(1)
 \end{aligned}$$

Bounding maximum degree $D(n)$

In previous section, we established that EXTRACT-MIN, and DELETE operations have amortized cost of $D(n)$, maximum degree of any node in Fibonacci heap. Lets now prove that $D(n) \leq O(\log n)$

Lemma 19.1 degree lower bound

Let x be any node in a Fibonacci heap, and suppose that $x.\text{degree} = k$. Let y_1, y_2, \dots, y_k denote the children of x in the order in which they were linked to x , from earliest to latest. Then, $y_1.\text{degree} \geq 0$ and $y_i.\text{degree} \geq i-2$ for $i = 2, 3, \dots, k$

→ Proof: By induction

• base case: attach y_1 to x

Start with a single node x , if we attach y_1 to x (through CONSOLIDATE), then $y_1.\text{degree} \geq 0$

• Inductive step: attach y_i to x for $i = 2, 3, \dots, k$

We know that when y_i attach to x , all previous y_1, y_2, \dots, y_{i-1} are already attached to x , and degrees must equals to attach, so:

$$y_i.\text{degree} \geq i-1$$

But, we also know that a node can lose at most one child before it becomes a root (if it lose 2 children, it gets cut from its parent), hence:

$$g_i \cdot \text{degree} \geq i - 2$$

Lemma 19.2 different expression Fibonacci number

Fibonacci numbers are defined by the recurrence:

$$F_k = \begin{cases} 0 & \text{if } k = 0 \\ 1 & \text{if } k = 1 \\ F_{k-1} + F_{k-2} & \text{if } k \geq 2 \end{cases}$$

But it can also be written as:

$$F_{k+2} = 1 + \sum_{i=0}^k F_i$$

Proof: By induction

• Base case: $k = 0$

$$\begin{aligned} F_2 &= 1 + \sum_{i=0}^0 F_i \\ &= 1 + F_0 \end{aligned}$$

• Inductive step: Assume $F_{k+1} = 1 + \sum_{i=0}^{k-1} F_i$, and we know that:

$$\begin{aligned} F_{k+2} &= F_{k+1} + F_k \\ &= 1 + \sum_{i=0}^{k-1} F_i + F_k \\ &= 1 + \sum_{i=0}^k F_i \end{aligned}$$

Lemma 19.3 Fibonacci number lower bound

For all integers $k \geq 0$, the $(k+2)^{\text{nd}}$ Fibonacci number satisfies $F_{k+2} \geq \phi^k$, where ϕ is the golden ratio, with value:

$$\phi = (1 + \sqrt{5})/2 = 1.61803\dots$$

Proof: By induction

• Base case: $k = 0$

$$\begin{aligned} F_{0+2} &= 1 + F_0 \\ &= 1 + 0 = \phi^0 \\ &= 1 \end{aligned}$$

$$k = 1$$

$$\begin{aligned} F_{1+2} &= 1 + F_0 + F_1 \\ &= 2 \geq \phi^1 \\ &= 1.61803\dots \end{aligned}$$

Inductive step: $k \geq 2$, Assume $F_i \geq \phi^{i-2}$ for $i = 0, 1, 2, \dots, k-1$

$$\begin{aligned} F_{k+2} &= F_{k+1} + F_k \\ &\geq \phi^{k-1} + \phi^{k-2} \\ &= \phi^{k-2} (\phi + 1) \\ &= \phi^{k-2} \cdot \phi^2 \\ &= \phi^k \end{aligned}$$

Lemma 19.4 Fibonacci heap node size lower bound

Let x be any node in a Fibonacci heap, and let $k = x.\text{degree}$

Then:

$$\text{size}(x) \geq F_{k+2} \geq \phi^k$$

$$\text{where } \phi = (1 + \sqrt{5})/2$$

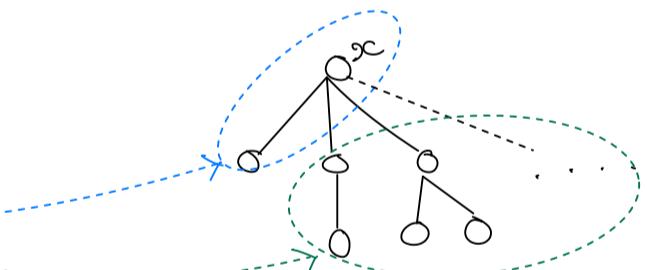
Proof:

- let s_k denotes minimum size of any node of degree k

$$\Rightarrow \begin{cases} s_0 = 1 \\ s_1 = 2 \end{cases}$$

node x_k is at most $\text{size}(x)$ since adding more node into x cannot decrease its size:

$$\begin{aligned} \text{size}(x) &\geq s_k \\ &\geq 2 + \sum_{i=0}^k s_{y_i.\text{degree}} \\ &\geq 2 + \sum_{i=0}^k s_{i-2} \quad \langle \text{lemma 19.1} \rangle \end{aligned}$$



- Prove $s_k \geq F_{k+2}$ by induction:

. Base case: $k=0$ and $k=1$

$$s_0 = 1 \geq F_2 = 1$$

$$s_1 = 2 \geq F_3 = \phi$$

- Inductive step: $k \geq 2$, assume $s_i \geq F_{i+2}$ for $i = 0, 1, \dots, k-1$

$$s_k \geq 2 + \sum_{i=0}^k s_{i-2}$$

$$\geq 2 + \sum_{i=0}^k F_i \quad \langle \text{hypothesis} \rangle$$

$$= 1 + \sum_{i=0}^k F_i$$

$$= F_{k+2} \quad \langle \text{lemma 19.2} \rangle$$

$$\geq \phi^k \quad \langle \text{lemma 19.3} \rangle$$

Corollary 19.5 Rank upper bound

The maximum degree $D(n)$ of any node in an n -node Fibonacci heap is $O(\lg n)$

Proof:

Let $\{n\}$ be number of nodes in Fibonacci heap
 x be any node in the same heap

Then by Lemma 19.4, we can say that:

$$n \geq \text{size}(x) \geq \phi^k$$

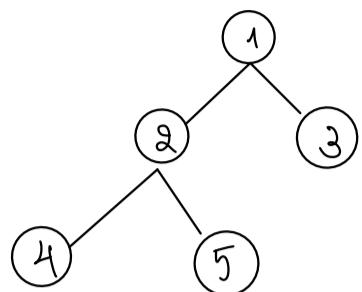
Taking log base ϕ gives us:

$$k \leq \log_{\phi} n$$

Therefore: The maximum degree $D(n)$ of any node is $O(\lg n)$

Reasoning through Fibonacci Heaps

We know ab binary heap



with 2 main properties:

1. Every level is full
2. Insert from left → right, top → bottom

And the analysis of its operations:

- GET_MIN : $O(1)$
- INSERT : $O(\lg n)$
- DECREASE_KEY : $O(\lg n)$
- EXTRACT_MIN : $O(\lg n)$ → important operation

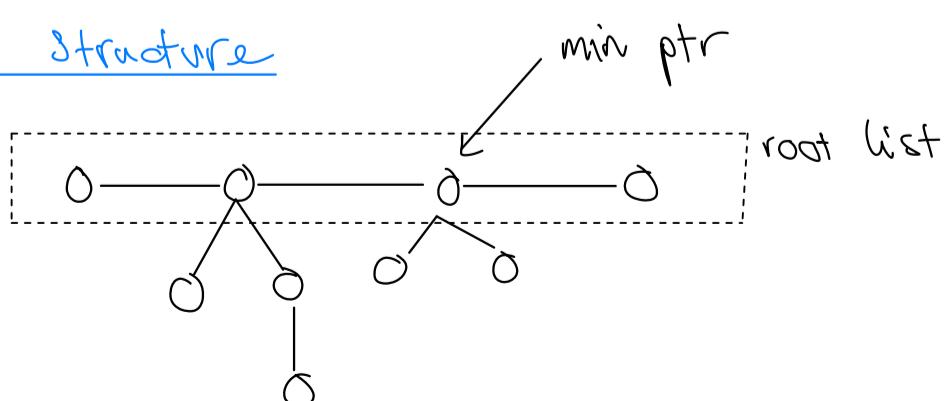
Can easily see that in order to EXTRACT_MIN, we need to maintain this structure that is "costly" for other operations. Can we do better?

↳ Yes, with Fibonacci Heap

Main idea of Fibonacci Heap is: "If we can somehow be lazy in all other operations and only do the work when it matters, i.e EXTRACT_MIN, then using Amortized analysis, we can argue that we have a data structure with constant time operations"

So, if-	GET_MIN	$O(1)$	}	\Rightarrow	Constant amortized cost
	INSERT	$O(1)$			
	DECREASE_KEY	$O(1)$			
	EXTRACT_MIN	$O(\lg n)$			

Overall Structure



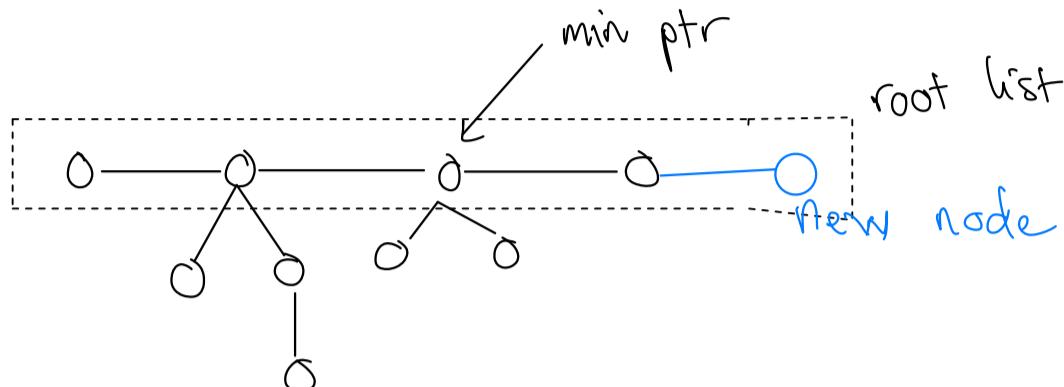
Maintain only 1 property: "Children must be larger than parent"

GET_MIN : $O(1)$

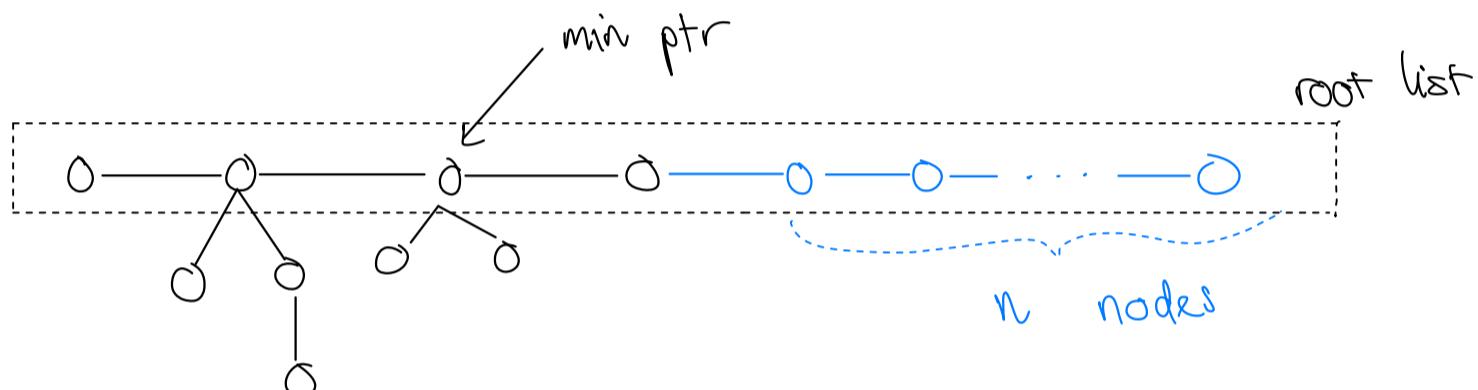
Because we maintain a "min-ptr"

INSERT : $O(1)$

"A new node is a new tree / root"



"If we insert n new nodes, then we get n new trees"



EXTRACT_MIN : $O(\lg n)$

This is where we "pay" for all that n nodes we INSERT
↳ Solution here is to somehow "merge" all n nodes into 1 now tree

An abstract example of a series of INSERT(s) and EXTRACT_MIN(s)

Operations	Costs	Amortized
INSERT	1	2
INSERT	1	2
⋮	⋮	⋮
INSERT	1	2
EXTRACT_MIN	100	10
EXTRACT_MIN	15	15
EXTRACT_MIN	15	15
⋮	⋮	⋮

Annotations:

- A bracket on the left side groups the first 100 operations (the first two rows plus the "⋮" row).
- The word "operations" is written next to the bracket.
- An arrow points from the "EXTRACT_MIN" row to the "Amortized" value "10", with the text "→ 1st EXTRACT_MIN operation".
- The word "pay" is written next to the "Amortized" value "15" in the second "EXTRACT_MIN" row, with the text "'pay' for previous INSERT".
- The word "s.t." is written next to the "Amortized" value "15" in the third "EXTRACT_MIN" row, with the text "the subsequent EXTRACT_MIN are constant".

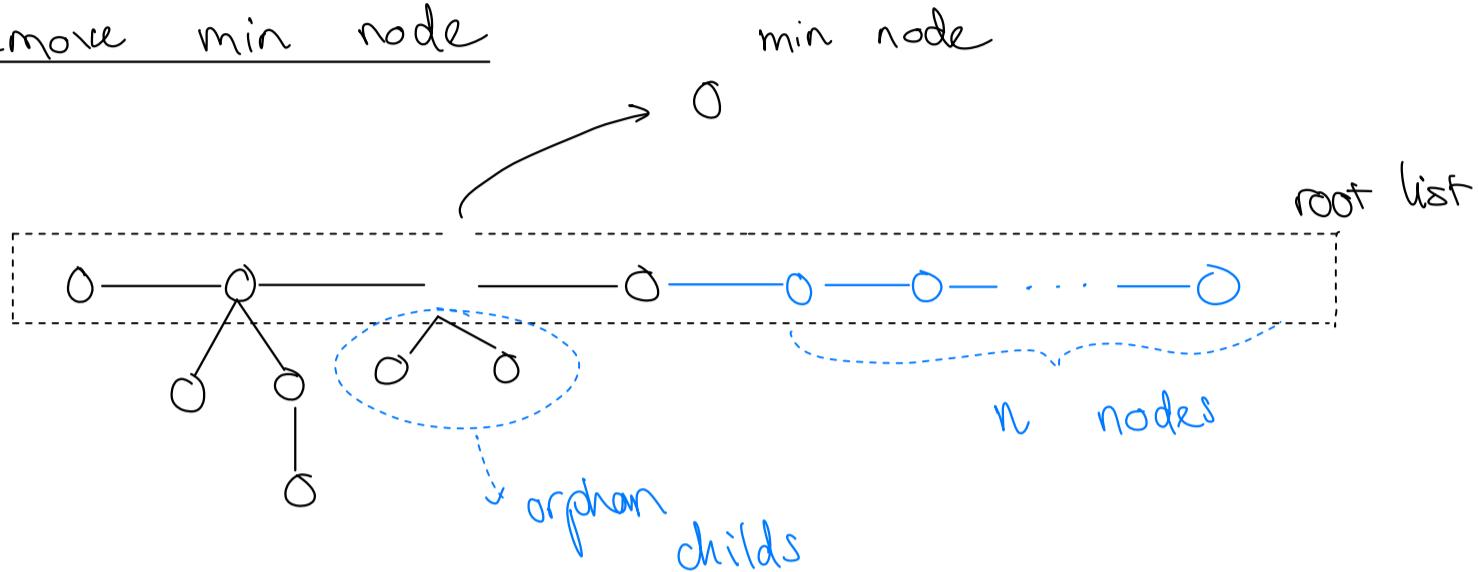
⇒ We need some EXTRA_WORK inside EXTRACT_MIN to allow the first EXTRACT_MIN to "pay" for previous INSERT(s)

\Rightarrow Amortized cost for EXTRACT-MIN = $O(\# \text{root after} + \text{EXTRA_WORK})$

\hookrightarrow The goal is that: "If we can come up with a EXTRA_WORK bound that is reasonably fast, then EXTRACT-MIN will be fast"

EXTRACT-MIN procedure

1. Remove min node

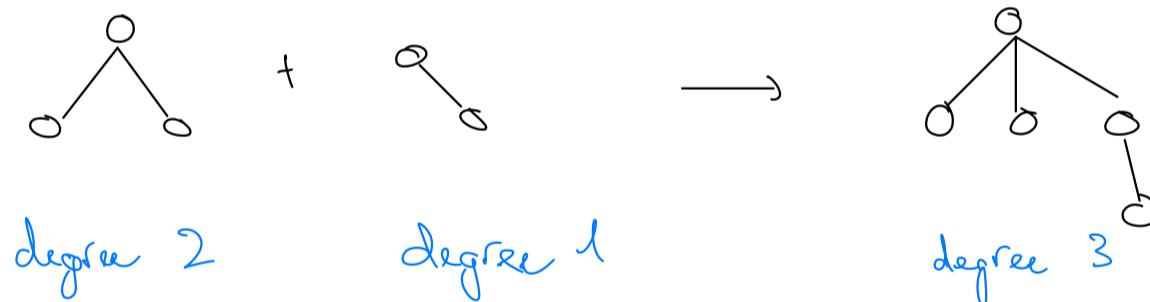


$\rightarrow O(\max \text{ degree})$ remove min node & set children as roots

2. CLEAN-UP

"Merge trees with the same degree until the root list only contains unique degree"

Merge:



Store in intermediate array:

0	1	2	3	...	Max degree

$\Rightarrow O(\max \text{ degree} + \# \text{roots})$: for each node in root list, "slide" it through intermediate array size Max-degree until find an empty spot.

3. Rebuild heap from intermediate array

Iterate through the array and build new heap

$\Rightarrow O(\max \text{ degree})$

Conclusion: EXTRACT-MIN cost is $O(\max\text{-degree} + \# \text{roots})$

Recall from the "abstract example", our amortized cost analyzed is $O(\# \text{roots} + \text{EXTRA-WORK})$

$$\Rightarrow O(\text{Max degree}) = O(\text{EXTRA-WORK})$$

↳ Let's bound Max degree (to show that EXTRACT-MIN is fast)

Informal proof:

Consider binomial trees, which has:

$$\text{degree } d \rightarrow 2^d \text{ nodes}$$

$$\Rightarrow k \text{ nodes} \rightarrow \text{degree } \log_2(k)$$

$$\Rightarrow \text{Max degree of tree with } n \text{ nodes is } \log_2(n)$$

So far: We have shown

GET-MIN

INSERT

EXTRACT-MIN

DECREASE-KEY

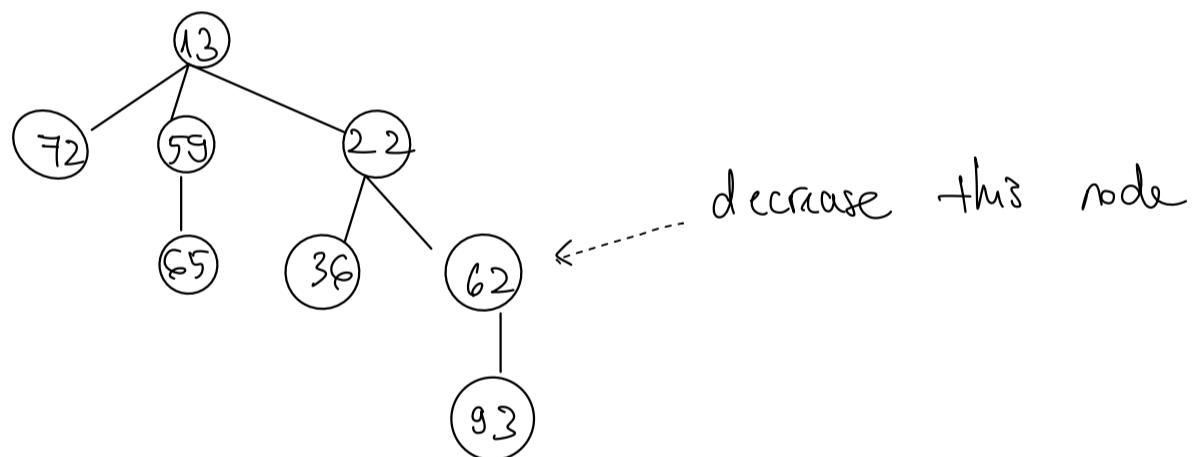
$O(1)$

$O(1)$

$O(\log n)$

?

DECREASE-KEY:



There are 2 scenarios:

Case 1: $x > x\text{-parent}$ i.e. decrease $62 \rightarrow 30$

\Rightarrow Do nothing

Case 2: $x < x\text{-parent}$ i.e. decrease $62 \rightarrow 15$

We do a few approaches

- Approach 1: Bubble up (like Binary heap)
 $\Rightarrow O(\log n)$

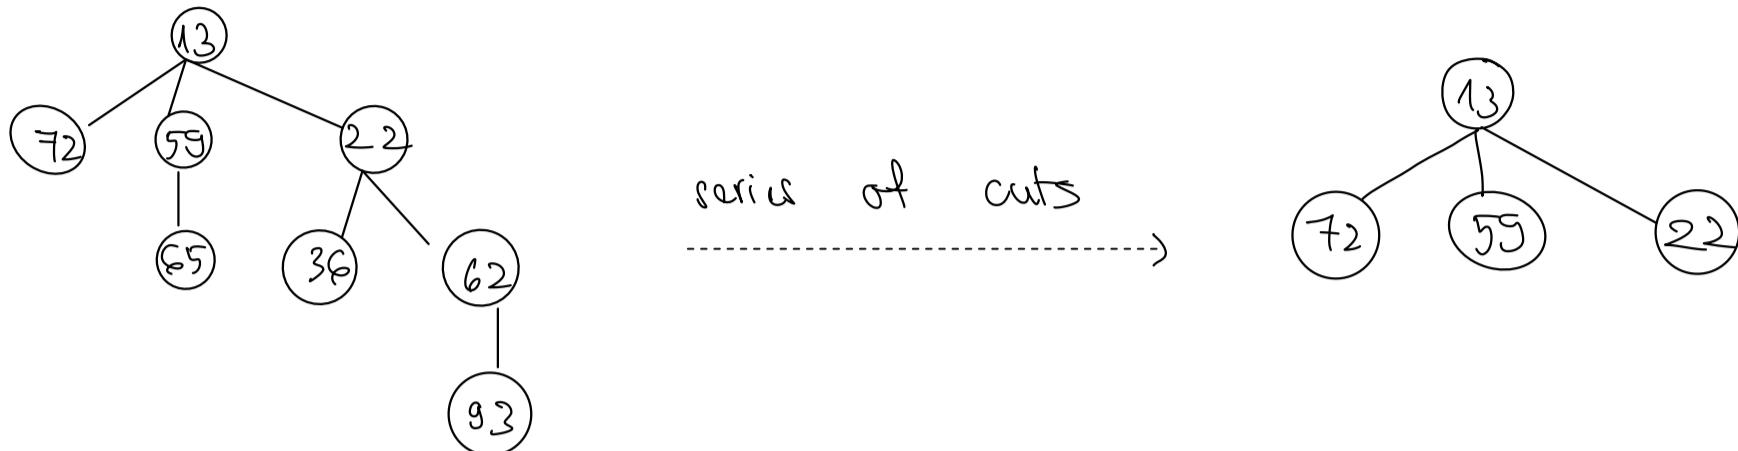
- Approach 2: Cut out $T(x)$ and form a new root
 $\Rightarrow O(1)$

\Rightarrow The goal is to be "lazy" and let EXTRACT-MIN clean up the mess, and end up with a good amortized cost (?)

↳ Problem with Approach 2:

If we only cut out the tree, then it is possible that we end up with a tree that is "high degree, but small # of nodes".

For example:



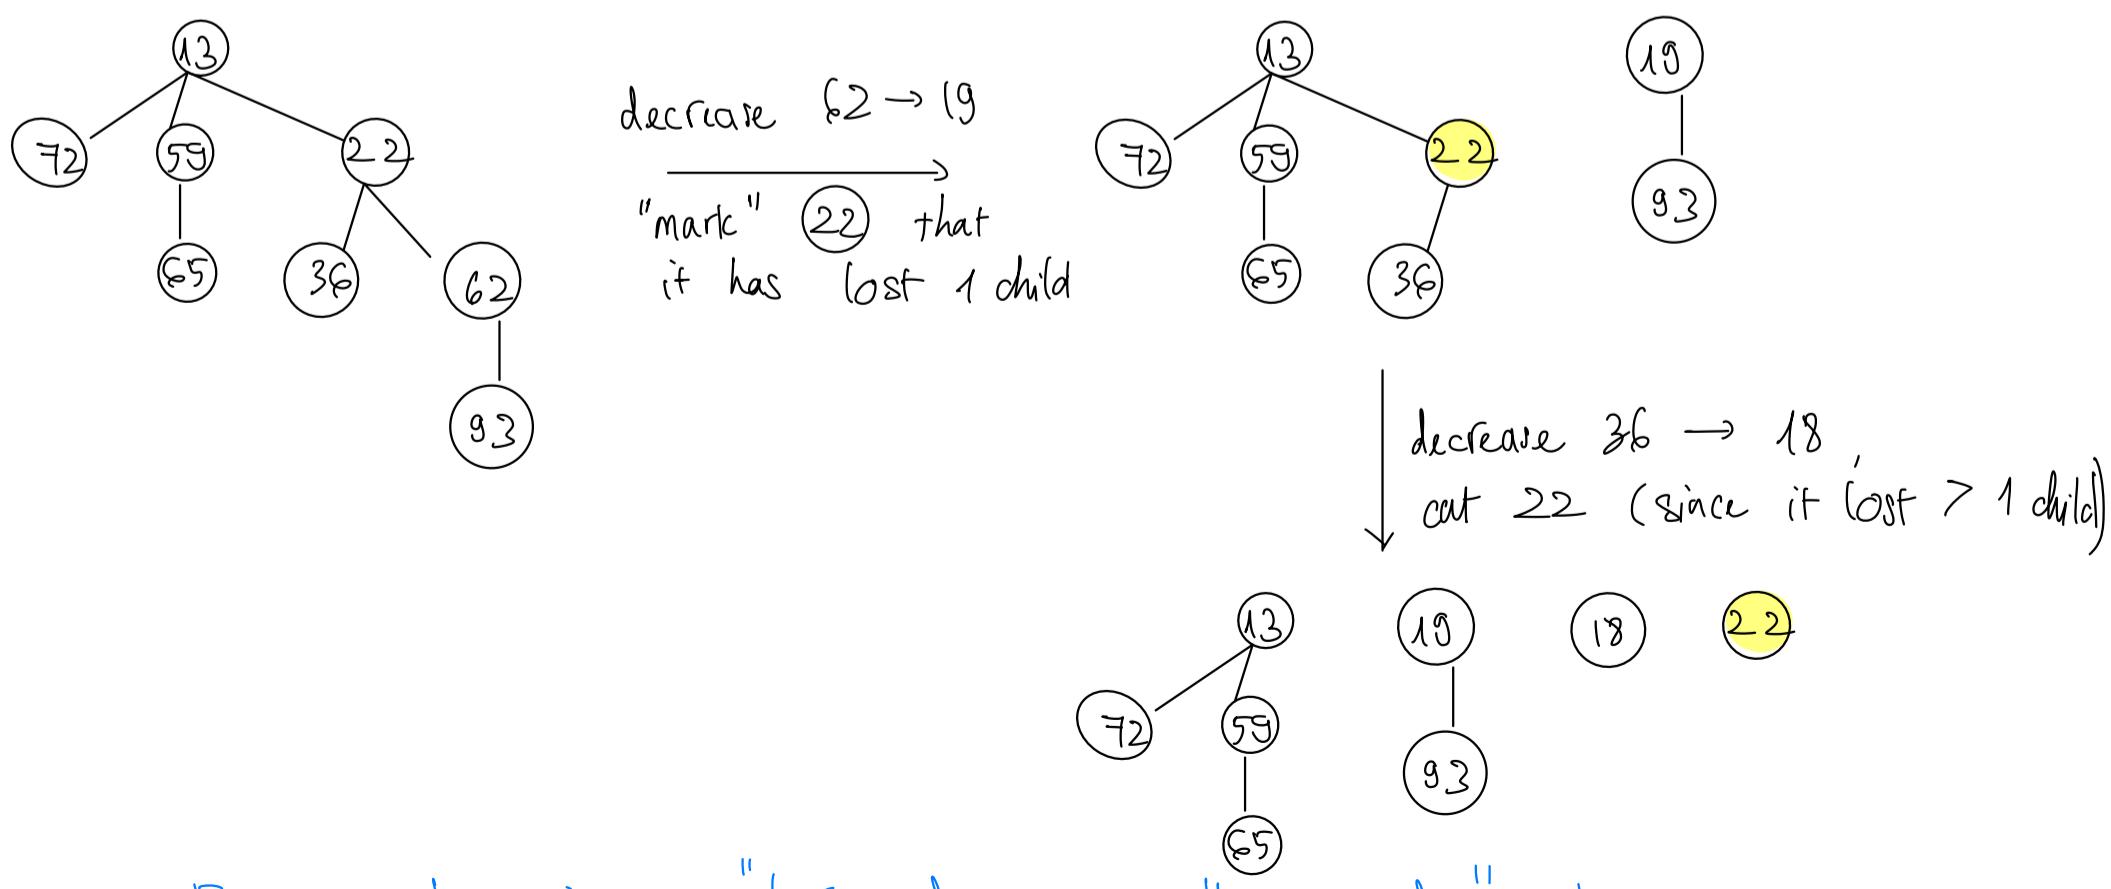
- binomial tree
- logarithmic degree
- not binomial tree
- linear degree

And recall that EXTRACT-MIN is bounded by $O(\text{max-degree})$
 \Rightarrow Degree grow linearly \Rightarrow EXTRACT-MIN cost $O(n)$, which
 destroy our proof that EXTRACT-MIN is "fast".

↳ Solution: Cut out at most 1 child per node

The goal is to come up with a procedure/rule that ensure we don't stray too much from the binomial tree shape
 \Rightarrow "If a node has more than one node cut out, that node is also cut out"

Example:



\Rightarrow Ensure there is no "large degree, small # nodes" tree

Proof : DECREASE_KEY is $O(1)$

- Question : If DECREASE_KEY is recursive , meaning we can DECREASE many parents on one call to DECREASE_KEY , how can we claim that DECREASE_KEY costs $O(1)$?
 - For k DECREASE_KEY called , cut out nodes $\leq 2k$
Because:
 - | . $\leq k$ nodes are cut on k DECREASE_KEY
 - | . $\leq k$ additional "marked" nodes are cut
- \Rightarrow Amortized cost of DECREASE_KEY is $O(1)$