

LOSS FUNCTION

Loss function for different models: (on a data point)

Linear Regression: $\hat{y} = wx + b$

• L2 loss (distance): $l(\theta | x, y) = \frac{1}{2} \| y - \hat{y} \|_2^2$

Binary classification: $\hat{y} = \sigma(wx + b)$

• Binary cross entropy: $l(\theta | x, y) = -y \log(\hat{y}) - (1-y) \log(1-\hat{y})$

Multi-class classification: $\hat{y} = \text{softmax}(wx + b)$

• Cross entropy loss: $l(\theta | x, y) = - \sum_{c=1}^C 1_{y=c} \log(\hat{y}_c)$

\Rightarrow basically $-\log(\hat{y}_y)$ \rightarrow groundtruth

Loss function of model is expected loss over dataset

$$L(\theta | D) = \mathbb{E}_{(x,y) \sim D} [l(\theta | x, y)]$$

Minimizing Loss function (Training model)

Problem statement: Find $\theta^* = \arg \min_{\theta} L(\theta | D)$

Use gradient descent: $\theta' = \theta - \epsilon [\nabla_{\theta} L(\theta)]^T$

ϵ \rightarrow learning rate
 $\nabla_{\theta} L(\theta)$ \rightarrow gradient

Initialization:

Usually random, according to Gaussian: $N(\mu, \sigma)$
or Uniform: $\text{Uniform}(a, b)$